

CS 5800 Distributed OS - Spring 2017

Homework 2

Josh Herman
Dalton Cole
Neil Blood

February 23, 2017

Problem 1 Suppose there are two GTAs (Alice and Bob) for CS384, each of whom is responsible for one section. In other words, Alice works as a grader for section 1 (S_1), and Bob works as a grader for section 2 (S_2). At the end of the semester, final grades will be assigned to students in the class based on their final total points. Assume there are two possible grades A and B , and Alice and Bob know the final total points for every student in their own sections (Alice does not know the final total points of students in S_2 , and vice versa).

You can assume that $|S_1| = |S_2| = n$ and $S_1 \cap S_2 = \emptyset$. Let $n = 3$, $S_1 = \{s_{11}(91); s_{12}(92); s_{13}(80)\}$ and $S_2 = \{s_{21}(81); s_{22}(82); s_{23}(97)\}$, where the value in the parentheses is the final total points for each student. At the end of your protocol, Alice will give A to s_{11} , s_{12} and B to s_{13} . Similarly, Bob will give A to s_{23} and B to s_{21} , s_{22} .

1. Design a distributed algorithm that assigns A to the upper half of all students and B to the lower half of all students in the class. Additionally, please make sure to minimize the communication cost of your protocol.

With the function $getMedian(A, B)$ Alice and Bob could figure out the overall median with a minimal communication cost and then both Alice and Bob could go through their section's grades and assign A 's to those above the overall median and B 's to those below the overall median.

```

getMedian(A, B)
{
    subA = A
    subB = B
    while(size(subA) > 2)
    {
        mA = median(subA)
        mB = median(subB)
        Alice send mA to Bob
        Bob send mB to Alice
        if mA == mB :
            return median is mA
        else if mA > mB
            subA = [subA[0]...mA]
            subB = [mB...subB[last]]
        else if mA < mB
            subA = [mA...subA[last]]
            subB = [subB[0]...mB]
    }
    Alice sends subA[0] and subA[1] to Bob
    Bob computes : median =  $\frac{\max(\text{subA}[0], \text{subB}[0]) + \min(\text{subA}[1], \text{subB}[1])}{2}$ 
    Bob sends median to Alice
    return median
}

```

2. Analyze the complexity of the algorithm in terms of the number of rounds, the size of the message for each round (assume 2 bytes are sufficient to store a final total point), and the local or individual computational costs.

The number of rounds for the algorithm, the while statement, would take $\log(n)$ steps and there would be two messages being sent each round, totaling 4 bytes. At the end of the while loop there would need to be two messages sent by Alice, totaling 4 bytes. So added all together the algorithm would cost $4 \log(n) + 4$ bytes. The computational cost would be $O(\log(n))$ since that's the number of times the while loop would iterate.

Problem 2 Define a well-founded set to prove that the program terminates in a bounded number of steps. You can assume that t is indeed the value of one of the elements in the array X .

The program is guaranteed to finish in n steps. This is because, in the worst case, $x[n-1] = t$. This is true because t is guaranteed to be in the set. In other words:

Set $s = \{\}$

The max size of s will be n . For every $x[i] \neq t$, append $x[i]$ to s . This will

represent all of the elements before t in the list. The max size of the set is n , which is the maximum number of steps to terminate the program.

Problem 3 Describe the desired safety and liveness properties when we use Google Drive to share files.

Safety: When two or more users edit the same document at the same time, their edits are never lost.

Liveness: Document will auto-save eventually, saving the data.

Problem 4 Refer to slide #19 of "handout4.pdf", prove a different example (from the one given on slide # 21) to prove that the program does not guarantee termination.

State	Action	c[0]	c[1]	c[2]	c[3]
A		0	0	2	0
B	P_0 moves	2	0	2	0
C	P_0 moves	0	0	2	0

As can be seen from the above table, states A and C are identical and the program returned to its initial state, thus the program can cycle with this given input, thus not guaranteeing termination.

Problem 5 Refer to slide #30 of "handout4.pdf", modify the program in a way that the safety property is still preserved, but the mathematical representation of the invariant is different from the one given on slide #31. Using induction to prove the safety property holds.

```

define  $c1, c2 : channel$ 
initially  $c1 = c2 = \emptyset$ 
{program for T}
defin $t : integer(initially\ t = 6)$ 
1   do  $t > 0$   $\rightarrow$  send message along  $c1; t := t - 1$ 
2    $\neg empty(c2)$   $\rightarrow$  receive message from  $c2; t := t + 1$ 
   od
{program for R}
defin $r : integer(initially\ r = 6)$ 
3   do  $\neg empty(c1)$   $\rightarrow$  receive message from  $c1; r := r + 1$ 
4    $r > 0$   $\rightarrow$  send message along  $c2; r := r - 1$ 
   od

```

Invariant changes to: $I \equiv (t \geq 0) \wedge (r \geq 0) \wedge (n1 + t + n2 + r = 12)$

Prove by induction that the safety property holds:

Base Case:

Initially: $n1 = n2 = 0, t = r = 6 \Rightarrow n1 + t + n2 + r = 12$

Induction Step: assume that I holds in the current state. We need to show that I holds after the execution of every eligible guarded action.

1. **After action 1:** The values of $(t + n1), n2, r$ remain unchanged. Also, since the guard is true when $t > 0$ and the action decrements t by 1, the condition $t \geq 0$ holds. Thus, I holds.
2. **After action 2:** The values of $(t + n2), n1, r$ remain unchanged. Also, since the value of t can only be incremented, so $t \geq 0$ holds. Thus, I holds.
3. **After action 3:** The values of $(r + n1), n2, t$ remain unchanged. Also, since the value of r can only be incremented, so $r \geq 0$ holds. Thus, I holds.
4. **After action 4:** The values of $(r + n2), n1, t$ remain unchanged. Also, since the guard is true when $r > 0$ and the action decrements r by 1, the condition $r \geq 0$ holds. Thus, I holds.

Therefore, with proof by induction the safety property still holds.

Problem 6 Show how to use the clock synchronization program (slide #30 of "handout4.pdf") to sync the clocks given in Figure 1.

2	0	1	0	2
1	2	2	2	1
0	0	0	0	0